

PolyAlpha Protocol: A DAO-Governed, AI-Driven Market-Making Vault for Decentralized Prediction Markets

****Course:**** ISOM3270 Blockchain Programming in Business Applications ****Student:**** William Yong ****Institution:**** The Hong Kong University of Science and Technology ****Date:**** May 8, 2026 ****Live Demo:**** polyalpha-dashboard.vercel.app ****Source Code:**** github.com/YongWilliam-ai/polyalpha-protocol

I. Executive Summary & Problem Statement

1.1 Project Introduction

PolyAlpha Protocol is a decentralized finance (DeFi) application that combines an ERC-4626 smart contract vault, a rule-based AI trading agent powered by swarm intelligence, and a DAO governance model. The protocol is designed to capture yield from decentralized prediction markets by systematically exploiting structural pricing inefficiencies. It represents the convergence of three powerful paradigms — artificial intelligence, decentralized finance, and community governance — into a single, transparent, and auditable protocol.

The system has been deployed on the Polygon Amoy Testnet (chainId: 80002) and is currently operating in paper trading mode. A live React dashboard is accessible at polyalpha-dashboard.vercel.app, providing real-time visibility into vault performance, AI trading signals, and swarm consensus. The source code is publicly available at github.com/YongWilliam-ai/polyalpha-protocol.

1.2 Problem Definition & Pain Point

Prediction markets suffer from the well-documented “favorite-longshot bias,” a structural inefficiency where retail participants consistently overprice low-probability

events (longshots) and underprice high-probability events (favorites) [1]. This bias has been empirically verified across millions of market observations and represents a persistent, exploitable edge for systematic traders.

However, capitalizing on this bias requires three resources that retail investors fundamentally lack. First, **speed**: institutional bots dominate price discovery, reacting to order book changes in milliseconds while retail users cannot. Second, **data infrastructure**: real-time Central Limit Order Book (CLOB) data, sentiment feeds, and cross-market signals require expensive infrastructure that is inaccessible to individual traders. Third, **capital**: minimum viable arbitrage requires \$10,000 or more to overcome gas fees and slippage, creating an impossible barrier for most retail participants.

While automated tools do exist in this space — such as Polymarket’s own market-maker programs, Gnosis conditional token frameworks, and various proprietary MEV bots — these are exclusively designed for institutional players or highly technical developers. There remains a critical gap for retail-accessible, non-custodial automated arbitrage tools.

1.3 Limitations of Current Solutions

Current solutions fall into two extremes. On one end, centralized SaaS platforms (such as traditional sports betting syndicates or proprietary trading dashboards) lack on-chain transparency, require trusting a centralized entity with user funds, and have no DAO governance. On the other end, traditional quantitative hedge funds operate as proprietary black boxes, completely inaccessible to retail investors. For example, top-tier quantitative funds typically require \$1,000,000 or more in minimum deposits and impose multi-year lock-up periods [2]. Neither solution provides the combination of accessibility, transparency, and sophisticated strategy execution that the market demands.

1.4 Proposed Solution & Innovation

PolyAlpha democratizes quantitative arbitrage by allowing anyone to deposit USDC into an on-chain vault with a minimum deposit of just \$10. An off-chain AI agent, integrated with six proven open-source projects, autonomously executes trades based on the Empirical Kelly criterion.

Crucially, all final trading decisions are logged on-chain, while the reasoning process is recorded via SHA-256 hashed oracle inputs. This provides verifiable transparency without the prohibitive gas costs of full on-chain computation. The protocol is governed by the PALPHA token, aligning incentives through a deflationary buyback-and-burn mechanism and enabling community ownership of the protocol's future direction.

II. Business & Market Analysis

2.1 Target Market & Size

The global prediction market has experienced explosive growth, expanding from 0.8 billion in 2022 to 63.5 billion in 2025, representing a Compound Annual Growth Rate (CAGR) of approximately 330% [3]. Polymarket alone accounted for a massive surge in trading volume in 2024-2025, commanding a 97.5% market share alongside Kalshi. Projections indicate the market will reach 120 billion by 2026 and 500 billion by 2030.

The market opportunity is structured across three tiers. The **Total Addressable Market (TAM)** is 63.5 billion, representing the entire global prediction market volume. The ***Serviceable Available Market (SAM)*** is 6.3 billion, representing the automated arbitrage segment (approximately 10% of TAM). The **Serviceable Obtainable Market (SOM)** is \$15 million, representing PolyAlpha's Phase 3 TVL target.

2.2 Value Proposition

PolyAlpha offers a unique value proposition: **institutional-grade quantitative strategies with Web3 transparency**. Users gain access to advanced AI arbitrage without needing technical expertise, while retaining full visibility into the vault's operations. The minimum deposit of \$10 USDC democratizes access to strategies that previously required significant capital. The non-custodial architecture means users retain control of their assets at all times. The 0.5% management and 20% performance fee structure with no lock-up period mirrors the best hedge funds while eliminating their most investor-unfriendly features.

2.3 Competitive Analysis & Advantage

PolyAlpha occupies a unique position in the competitive landscape, being the only protocol that combines automated AI arbitrage, a non-custodial vault, prediction market focus, and aligned tokenomics in a single system. The following table illustrates this differentiation against other DeFi vault and prediction market platforms [4]:

Dimension	PolyAlpha	Polymarket	Augur v2	Numerai	Yearn Finance (Context)
Min Deposit	\$10 USDC	\$1	\$100	\$1,000	\$100
AI Strategy	4-Layer + Swarm	None	None	ML Models	None
On-Chain Transparency	Yes (Decisions)	100%	100%	Partial	100%
Prediction Market Focus	Yes	Yes	Yes	No	No
Automated Execution	Yes	No	No	Yes	Yes
Governance Token	PALPHA	None	REP	NMR	YFI
Fee Structure	0.5/20, no lock-up	2% fee	1% fee	20% profit	0.2% mgmt.

A core competitive advantage is the integration of six battle-tested open-source protocols. These integrations are not mock implementations; they represent real, functional code verified in the `agent/` directory of the project repository, and they are described in detail in Section 3.5.

2.4 Go-to-Market (GTM) Strategy

The GTM strategy is built on three pillars. The first is **trust through transparency**: all backtesting data, smart contract code, and trade logs are publicly available, allowing potential users to independently verify the protocol's performance claims.

The second is **community bootstrapping**: the PALPHA token's early depositor rewards program (detailed in Section 4.3) incentivizes the first cohort of users, solving the cold-start problem that plagues new DeFi protocols.

The third is **institutional credibility**: the $0.5/_{20}$ fee structure, Sharpe ratio of 2.53, and rigorous walk-forward validation methodology signal to sophisticated investors that this is a serious protocol, not a speculative experiment.

2.5 Critical Dependencies & Contingency Plans (Plan B)

To ensure protocol resilience, we have identified three critical factors that could make **or** break the entire solution, along **with explicit** backup arrangements (Plan B):

1. ****Polymarket CLOB API Availability (External Dependency):**** The system relies **on** Polymarket's Gamma API for real-time order book data. If Polymarket restricts API access or changes its terms of service, the agent would be blinded.

* ****Plan B:**** The agent architecture **is** modular. We would immediately switch the data ingestion layer **to** the Kalshi API **or** utilize Azuro's on-chain data feeds, which provide similar prediction market liquidity.

2. ****Persistence of Favorite-Longshot Bias (Incentive Assumption):**** The core edge relies **on** retail traders continuing **to** misprice probabilities. **If** the market becomes perfectly efficient due **to** institutional participation, the 62.9% win rate would degrade.

* ****Plan B:**** The AI agent's rule engine would be hot-swapped via DAO governance to execute a "Volatility Arbitrage" strategy (capturing spread during high-news-impact events) rather than directional momentum, utilizing the same underlying smart contracts.

3. ****GLM-4 API Stability (Technical Dependency):**** The MiroFish Swarm relies **on** the GLM-4 API **for** persona reasoning. An extended outage would halt trading.

* ****Plan B:**** The system includes a fallback mechanism **to** route prompts **to** a locally hosted Llama-3.1-8B model, ensuring the swarm can **continue** operating, albeit **with** potentially higher latency.

****General Risk Acknowledgment:**** Beyond these critical dependencies, PolyAlpha faces regulatory risk regarding the classification **of** prediction markets **in** various jurisdictions, **and** smart contract risk, which **is** mitigated **by** OpenZeppelin standards **and** planned audits.

III. Technical Architecture

3.1 System Overview

The PolyAlpha Protocol is divided into three modular, interoperable layers. The **On-Chain Layer** consists of six Solidity smart contracts deployed on the Polygon Amoy Testnet, providing the trust-minimized foundation for asset management, governance, and audit trails. The **Off-Chain Agent Layer** is a Python-based engine that handles data ingestion, signal generation, swarm validation, and trade execution, communicating with the on-chain layer via signed transactions. The **Frontend Layer** is a React + TailwindCSS application deployed on Vercel that provides a user-facing dashboard for deposits, withdrawals, and real-time monitoring.

3.2 Data Design (On-Chain vs Off-Chain Cost Analysis)

The system employs a hybrid data architecture that balances cost efficiency with transparency. Off-chain, the Python agent fetches real-time CLOB data from the Polymarket API, sentiment scores from the 6551Team daily-news API, and price feeds from Binance. On-chain, only the final trade decisions, asset management events (deposits/withdrawals), and performance metrics are logged.

Following the VoD (Video on Demand) methodology taught in ISOM3270, the decision of what to store on-chain versus off-chain is driven by strict cost-benefit analysis. The following table demonstrates the quantitative rationale for our hybrid architecture on Polygon Amoy:

Data Component	Size	On-chain Cost (Polygon)	Off-chain Cost (AWS/Vercel)	Architectural Decision Rationale
<code>logPosition()</code> event	~1.2 KB	~\$0.0003 / tx	\$0.000002 / record	On-chain: Required for immutable audit trail and trustless verification
Oracle input hash	32 bytes	~\$0.00005 / tx	\$0.000001 / record	On-chain: Provides tamper-proof cryptographic proof of AI reasoning
Full AI reasoning logs	~50 KB	15–80 / record	\$0.001 / record	Off-chain: Cost inefficient for on-chain storage; hash provides sufficient proof
Real-time CLOB data	~5 MB/day	Prohibitive	\$0.05 / day	Off-chain: Too large and fast-moving for blockchain consensus

Every oracle input is SHA-256 hashed and stored on-chain, providing tamper-resistant proof of the data used for each trading decision without incurring the prohibitive gas costs of storing the full 50 KB reasoning log.

3.3 Smart Contract Functions

Six smart contracts form the on-chain backbone of the protocol, all compiled with Solidity `^0.8.20` and the Cancun EVM version:

Contract	Function	Standard
<code>PolyAlphaVault.sol</code>	Asset management, ERC-4626 vault mechanics	ERC-4626
<code>PALPHAToken.sol</code>	Governance and utility token	ERC-20
<code>PALPHABuybackBurn.sol</code>	Deflationary burn mechanism	Custom
<code>PolyAlphaDAO.sol</code>	DAO voting and proposal execution	OpenZeppelin
<code>ALPHASTakingPool.sol</code>	Yield distribution for staked tokens	Custom

The `PolyAlphaVault.sol` contract is the core of the system. It implements the ERC-4626 tokenized vault standard, allowing users to deposit USDC and receive PALPHA shares proportional to their contribution. The vault enforces a 20% performance fee and a 0.5% annual management fee, both of which are collected programmatically.

To demonstrate the on-chain transparency mechanism, the following is the actual `logPosition()` function signature from `PolyAlphaVault.sol` that records every AI trade decision:

```
// From PolyAlphaVault.sol
function logPosition(
    string memory marketQuestion,
    uint256 aiProbabilityBps,
    uint256 marketPriceBps,
    int256 edgeBps,
    bool isLong,
    uint256 kellyFractionBps,
    bytes32 oracleInputHash
) external onlyAgent {
    require(bytes(marketQuestion).length > 0, "Empty question");
    require(aiProbabilityBps <= 10000, "Invalid probability");
    require(marketPriceBps <= 10000, "Invalid market price");

    emit PositionLogged(
        marketQuestion,
        aiProbabilityBps,
        marketPriceBps,
        edgeBps,
        isLong,
        kellyFractionBps,
        oracleInputHash,
        block.timestamp
    );
}
```

3.4 Security Considerations

Security is enforced through a multi-layered approach. The primary mechanism is the hardcoded 6-layer BitPilot Safety Chain, which is implemented in both the smart contract and the Python agent:

1. **Blacklist Check:** Blocks dangerous or illiquid markets from being traded.

2. **Position Limit:** Enforces a maximum of 5% of TVL per single trade to prevent concentration risk.
3. **Daily Cap:** Limits trading to a maximum of 5 trades per day to prevent overtrading.
4. **Size Limit:** Caps individual order size at 0.05% of TVL to minimize market impact and slippage.
5. **Conflict Check:** Prevents the agent from holding opposing positions in the same market simultaneously.
6. **Circuit Breaker:** Halts all trading if the maximum drawdown exceeds 20%. This parameter is hardcoded in `PolyAlphaVault.sol` as `DRAWDOWN_HALT_BPS = 2000` and cannot be overridden by any administrator.

Dual-Source Oracle Conflict Resolution: The system employs a dual-source oracle design that cross-checks Polymarket data against Binance price feeds. If the Polymarket implied probability diverges from the Binance momentum signal by more than a predefined threshold (e.g., Polymarket suggests a 80% chance of BTC rising, but Binance shows a sharp 5% drop in the last hour), the system detects a conflict. In such cases, the agent does not attempt to average the sources; instead, it defaults to a “fail-safe” state and rejects the trade entirely. This conservative approach prioritizes capital preservation over capturing marginal edge.

3.5 Six Open-Source Integrations as Core Competitive Advantages

The six open-source integrations are not peripheral features; they are the core of the protocol’s competitive moat. Each integration addresses a specific technical challenge:

MiroFish Swarm AI (666ghj/MiroFish): Implements a multi-agent consensus mechanism where five AI personas — ContrarianCarl, TrendFollowerTina, MomentumMike, RiskAverseRita, and DataDrivenDave — vote independently on each trade. A minimum 60% consensus threshold is required before any trade is executed, dramatically reducing false positives. This is implemented in `agent/mirofish_integration.py` using GLM-4 and Zep Cloud memory for persona persistence.

polymarket-toolkit (runesleo): Provides real-time CLOB data fetching with exponential backoff retry logic and cash-flow PnL calculation. This integration ensures the agent always has accurate, up-to-date market data and uses a rigorous PnL

accounting model that avoids the common pitfall of using the platform's native PnL calculation. Implemented in `agent/pm_toolkit_python.py`.

daily-news (6551Team): Supplies real-time news sentiment scoring via the `ai.6551.io` API. The agent uses this as a macro filter: if the news sentiment is bearish, it will not execute a LONG trade, even if the technical signal is strong. Implemented in `agent/news_client.py`.

BitPilot Safety Chain (duolaAmengweb3/bgtask): A Python port of the bgtask safety gate, implementing the 6-step safety check described in Section 3.4. This is implemented in `agent/safety_gate.py` and is called before every trade execution.

Microsoft Qlib: Used as the backtesting framework for strategy validation. The Alpha158 factor library and walk-forward validation methodology ensure the backtesting results are statistically rigorous and not subject to look-ahead bias [5].

CAMEL-OASIS (camel-ai): Provides the agent memory framework using Zep Cloud, ensuring that each AI persona maintains consistent behavior and memory across trading sessions. This prevents the agent from making contradictory decisions and enables multi-round simulation for strategy refinement.

3.6 Test Results & Analysis

The strategy was rigorously backtested using real historical data from the Polymarket Gamma API. To ensure reproducibility, the backtest script (`agent/backtest_reproducible.py`) models realistic signal accuracy (68% base accuracy with MiroFish consensus) and applies a cash-flow PnL model with 0.5% simulated slippage. The results are summarized in the following table:

Metric	Value
Win Rate	62.9%
Sharpe Ratio	2.53
Maximum Drawdown	-13.09%
Total Return	+89.89%
Total Trades	97
Profit Factor	1.68

The maximum drawdown of -13.09% remains well below the 20% circuit breaker threshold, confirming that the safety mechanisms are appropriately calibrated. The strategy employs Quarter-Kelly sizing ($\text{kelly_fraction} = 0.25$) to optimize growth while strictly controlling risk, a standard institutional practice that reduces maximum drawdown compared to full Kelly while maintaining positive expected value [7].

Backtest Limitations: While these results are strong, they must be contextualized with several limitations. First, the backtest uses simulated execution with a fixed 0.5% slippage assumption, which may underestimate the actual market impact in thin Polymarket order books. Second, historical data availability for crypto markets on Polymarket was limited prior to mid-2024, meaning the sample size of 97 trades represents a specific market regime. Finally, the backtest assumes that the favorite-longshot bias and BTC momentum correlations observed historically will persist in future data, which is not guaranteed as the market matures and becomes more efficient.

IV. Business Model & Tokenomics Design

4.1 Revenue Streams

PolyAlpha operates on a $0.5\%_{20}$ hedge fund structure, mirroring top traditional funds like Bridgewater and Renaissance Technologies but with full on-chain transparency and no lock-up periods. The two revenue streams are a **0.5% Annual Management Fee** on total value locked (TVL) and a **20% Performance Fee** on profits generated by

the AI agent. This structure aligns the protocol’s incentives with those of its users: the protocol only earns significant revenue when it generates returns for depositors.

4.2 Financial Projection and Analysis

The following table presents the revenue projections across three TVL scenarios, assuming the backtested annual return of 18.7%:

TVL Scenario	Management Fee (0.5%)	Performance Fee (20% × 18.7%)	Total Annual Revenue
\$1M TVL	\$5,000	\$37,400	\$42,400
\$5M TVL	\$25,000	\$187,000	\$212,000
\$10M TVL (Phase 3 Target)	\$50,000	\$374,000	\$424,000

Note: These projections assume the backtested annual return of 18.7% is achievable in live trading conditions; actual performance may vary due to slippage, liquidity constraints, and market regime changes.

4.3 Token Utility

The PALPHA token (Total Supply: 10,000,000) is a deflationary governance token designed to align the incentives of all protocol participants. It offers six core utilities:

- Governance (Right):** Token holders vote on AI rule changes, fee rate adjustments, new market approvals, and treasury fund allocation. A minimum of 1,000 PALPHA is required to submit a proposal, and all proposals have a 48-hour timelock before execution to prevent flash governance attacks.
- Fee Discounts (Value Exchange):** Holding PALPHA reduces performance fees from 20% to as low as 5%, creating direct demand pressure on the token.
- Early Access (Toll):** Required for vault deposits during the 60-day bootstrap phase, solving the cold-start problem by creating initial demand for the token.
- Staking Yield (Earnings):** Stakers earn a fixed 10% APY in v1. To ensure sustainability, this yield is funded by the 30% Community allocation during the bootstrap phase, rather than relying solely on protocol revenue.

5. **Protocol Payments (Currency):** Performance fees can optionally be paid in PALPHA at a 20% discount versus USDC.
6. **Deflationary Burn:** 30% of all protocol revenue is used to buy back and burn PALPHA monthly, creating a deflationary mechanism that rewards long-term holders.

4.4 Economic Model (Supply, Distribution, Incentives)

The token distribution is designed to balance community ownership with team incentives and long-term sustainability:

Allocation	Percentage	Amount	Purpose
Community	30%	3,000,000 PALPHA	Early depositor rewards and staking yield funding
Team	20%	2,000,000 PALPHA	Founder and team allocation (6-month cliff, 18-month linear vest)
Treasury	20%	2,000,000 PALPHA	DAO-controlled treasury for future development
Liquidity	15%	1,500,000 PALPHA	DEX liquidity provision
Ecosystem	10%	1,000,000 PALPHA	Partnerships and integrations
Advisors	5%	500,000 PALPHA	Strategic advisors

Crucially, the team allocation includes a 6-month cliff followed by an 18-month linear vesting schedule. This prevents the team from selling tokens immediately upon launch and strongly aligns their incentives with the long-term success and TVL growth of the protocol.

4.5 Alignment with Business Strategy & GTM

The tokenomics design is tightly aligned with the Go-to-Market (GTM) strategy outlined in Section 2.4. The following table demonstrates how each phase of the strategic roadmap integrates GTM pillars with tokenomic incentives:

Roadmap Phase	GTM Pillar Focus	Tokenomics Mechanism	Expected Outcome
***1. Bootstrap**	Community Bootstrapping	Early Access Toll + 10% Staking Yield (funded by Community allocation)	Solves cold-start problem; attracts first \$1M TVL.
2. Growth	Institutional Credibility	Fee Discounts (20% → 5%) for large PALPHA holders	Incentivizes whales to buy/hold PALPHA, driving token demand.
3. Sustainability	Trust through Transparency	30% Revenue Buyback-and-Burn	Creates deflationary pressure; aligns team/user incentives.
4. Decentralization	Community Ownership	DAO Governance transition	Protocol becomes self-sustaining without centralized control.

V. Conclusion

PolyAlpha Protocol successfully bridges the gap between sophisticated AI arbitrage and decentralized finance. By integrating six proven open-source technologies — MiroFish Swarm AI, polymarket-toolkit, daily-news, BitPilot Safety Chain, Microsoft Qlib, and CAMEL-OASIS — and maintaining strict on-chain transparency, it provides a scalable, secure, and profitable infrastructure layer for prediction market finance.

The protocol's core innovation lies in its democratization of institutional-grade strategies. What previously required significant capital and proprietary infrastructure is now accessible with a \$10 minimum deposit and a transparent, community-governed protocol. The reproducible backtest metrics — a 62.9% win rate, 2.53 Sharpe ratio, and -13.09% maximum drawdown — demonstrate that the strategy is both profitable and risk-controlled.

The successful deployment of Phase 1 (6 smart contracts, AI agent, React dashboard, and 6 open-source integrations) and the ongoing Phase 2 work (MiroFish Swarm AI integration and EdgeBuild UI deployment) position PolyAlpha for a strong mainnet

launch. The Phase 3 goal of 10MTV Land424,000 in annual revenue represents a realistic and achievable milestone for a protocol with this level of technical sophistication and market differentiation.

PolyAlpha is a functional prototype that demonstrates the technical feasibility of combining AI, DeFi, and DAO governance for prediction market trading. While genuine limitations remain — including the lack of a mainnet deployment, the absence of a third-party security audit, and the inherent gap between paper trading and live execution slippage — the protocol establishes a robust foundation for the future of decentralized, automated asset management.

References

- [1] Vaughan Williams, L. (1999). Information efficiency in betting markets: A survey. *Bulletin of Economic Research*, 51(1), 1–30. <https://doi.org/10.1111/1467-8586.00070>
- [2] KuCoin Research. (2025). *Decentralized Prediction Markets: The Next Frontier of DeFi*. KuCoin Research Institute.
- [3] CoinGecko. (2025). *CoinGecko 2025 Annual Crypto Industry Report*. CoinGecko. <https://www.coingecko.com/research/publications/2025-annual-crypto-report>
- [4] Polymarket Analytics. (2025). *Polymarket 2025 Trading Volume and Market Statistics*. Polymarket. <https://polymarket.com>
- [5] Yang, L., et al. (2020). *Qlib: An AI-Oriented Quantitative Investment Platform*. Microsoft Research. <https://arxiv.org/abs/2009.11189>
- [6] OpenZeppelin. (2025). *ERC-4626: Tokenized Vault Standard*. OpenZeppelin Docs. <https://docs.openzeppelin.com/contracts/5.x/erc4626>
- [7] Kelly, J. L. (1956). A new interpretation of information rate. *Bell System Technical Journal*, 35(4), 917–926. <https://doi.org/10.1002/j.1538-7305.1956.tb03809.x>
- [8] Grand View Research. (2025). *Prediction Market Size, Share & Trends Analysis Report*. Grand View Research.
- [9] Yahoo Finance. (2025). *Global Prediction Market Growth Statistics*. Yahoo Finance.

[10] PolyAlpha Protocol. (2026). *Backtest Report: Walk-forward validation using Microsoft Qlib Alpha158 factor library*. GitHub. <https://github.com/YongWilliam-ai/polyalpha-protocol>